TRAIL OFBITS

# Differential analysis of x86-64 decoders William Woodruff, Niki Carroll, Sebastiaan Peters

## **Contact Information**



# William Woodruff Trail of Bits

william@trailofbits.com @8x5clPW2

IEEE S&P WORKSHOPS | LANGSEC 2021



**Niki Carroll**George Mason University

ncarrol5@gmu.edu @inventednight



#### **Sebastiaan Peters**

Eindhoven University of Technology

s.peters2@student.tue.nl



#### Items

- Overview of instruction decoding
  - Why is x86-64 especially challenging?
- Instruction decoding as an attack surface
- Differential analysis of instruction decoders
- Mishegos: a differential fuzzer for x86-64 decoders
- Results and future work

## Instruction decoding

- Individual machine code instructions → assembler mnemonics
- The fundamental building block for correct disassembly
  - Function, control/data flow, call graph recovery all depend on correct decoding
  - Small mistakes → incorrect control/data flow, misaligned decodings, ...

#### Decoding depends on ISA/vendor-specific parameters:

- Fixed length or variable length encoding?
- Revisions to the ISA? Backwards compatible, incompatible?
- Open, machine-readable specification?
- Hidden opcodes/unintended functionality?
- Vendor-specific behavior?
- Completeness vs. decoding what the common compilers emit?

## Instruction decoding: not just for REs anymore!

#### Previously: primarily reverse engineers and debuggers

- Decoder errors here are annoying, but there's a human in the loop
- Human fixes/stubs the error and continues

#### Now: on the hotpath of user code:

- Interpreted languages: JITs, sandboxes
- Antivirus tools: static disassembly and analysis
- Static & dynamic binary translation (Apple's Rosetta)
- Self-modifying code (runtime patching in the Linux kernel)
- Constrained runtimes/VMs (WASM, eBPF)
- No human in the loop to catch errors!

## Can we exploit instruction decoding errors?

- Errors in instruction decoding have security consequences
  - Sandboxes, VMs, JITs: run mis-generated code, escape the environment
    - eBPF: violate runtime guards
  - Antivirus: confuse AV checks into skipping/not flagging malicious code
    - ...or trip up code/data disambiguation
  - SBTs/DBTs: convert a benign program into a misbehaving one
    - Turn benign code into vulnerable code
    - Introduce timing channels or information leaks?
- Speculation: x86-64's complexity makes it particularly susceptible to decoder implementation errors

## x86-64 decoding is (particularly) hard

#### Numerous features that make x86-64 hard to decode correctly:

- Complex, variable-length instruction format (up to 15 bytes per instruction)
- >40 years of semi-compatible ISA revisions
- No formal specification, two separate major vendors with references aimed at developers rather than decoder implementers
- A long history of undocumented opcodes, including backdoors!
  - intel: AAD/AAM variants, SALC, ICEBP, UD1
  - (2) AMD: 3DNow! variants (Domas 2017)
  - VIA: ALTINST + C3 family backdoors (Domas 2017)
- Variations in vendor behavior (operand size prefix, SYSRET differences)
- Under- and undefined semantics
  - Multiple prefixes, EFLAGS state on BSF/BSR variants

## Where are the bugs in x86-64 decoders?

- Speculation 1: many bugs will exist where discrete regions of the instruction format meet
  - Unlikely to completely mess up displacement decoding, but might mess up displacement decoding with multiple legacy prefixes and an uncommon SIB encoding
- Speculation 2: many bugs will also exist where x86-64 has undergone historical changes
  - Subtleties with the REX prefix, validity of instructions in long mode, ...
- Speculation 3: x86-64 instruction decoders are primarily tested against compilers
  - o Decoder authors under-test against instructions/formats that compilers don't use
- Speculation 4: most bugs in x86-64 instruction decoders won't cause memory unsafety in the decoder itself
  - Not a good target for traditional crash-driven fuzzing...

## Differential analysis of instruction decoders

Instruction decoders are an ideal target for differential fuzzing:

- Relatively little memory unsafety → simple harnesses
- Lots of competing (and open source!) implementations
- Simple interfaces (byte string → instruction)
- Lots of signals to diff against:
  - Decoding success/failure, failure kind
  - Decoded length
  - Decoded semantics (instruction ground, operand count, ...)

## Differential analysis of instruction decoders

#### **Prior work:**

- Paleari et al. (2010): Compare N decoders against a hardware decoder as a source of ground truth
- Jay & Miller (2018): Normalize each decoder's assembly output and compare for discrepancies

Both approaches focus on false negatives; we also want to identify false positives!

## Mutation strategies for x86-64

#### **Prior work:**

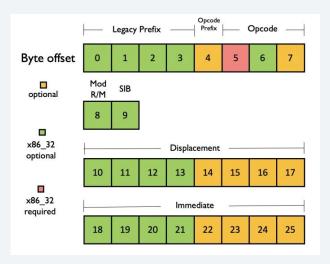
- Paleari: Generate random input sequences, mix in with sequences from hardware-guided generation
- Jay & Miller: Seed the mutator with valid inputs, perform bitflips and sample outputs based on inferred instruction structure

We want to do better than random generation without relying on hardware ground truth (Paleari) or structure inference from a potentially unreliable decoder (Jay & Miller)!

## A "sliding" mutation strategy

Observation: the maximum x86-64 instruction length is 15 bytes

... but the maximum size of each instruction field adds up to 26:



We can soundly overapproximate the structure of a potential x86-64 instruction by filling up to 26 bytes!

We can evaluate speculations 1 and 2 by being better-than-random about legacy prefixes, REX, ModR/M, SIB, etc.

## A "sliding" mutation strategy

#### Mutation engine lifecycle

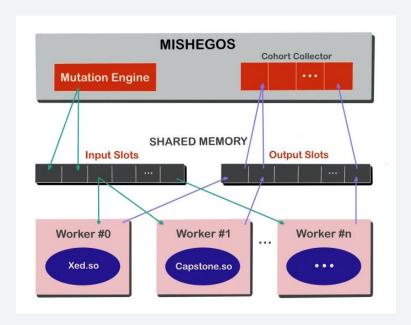
- Generate a "maximal" overapproximated candidate
- Extract individual 15-byte "sliding" candidates from the "maximal" candidate
- Once the "maximal" is exhausted, generate a new one
- Repeat until fuzzing is halted

```
if (insn_cand.off == 0) {
   build_sliding_candidate();
}

if (insn_cand.len <= MISHEGOS_INSN_MAXLEN) {
   memcpy(slot->raw_insn, insn_cand.insn, insn_cand.len);
   slot->len = insn_cand.len;
   insn_cand.off = 0;
} else {
   memcpy(
      slot->raw_insn,
      insn_cand.insn + insn_cand.off,
      MISHEGOS_INSN_MAXLEN);
   slot->len = MISHEGOS_INSN_MAXLEN;
   insn_cand.off =
      (insn_cand.off + 1) % (insn_cand.len - MISHEGOS_INSN_MAXLEN + 1);
}
```

Tests our speculations by feeding a wide range of prefix/structured interactions to the decoders!

#### Mishegos: our differential fuzzer for x86-64



- Each "sliding" candidate is placed in an input slot, to be consumed by each decoder's worker
- Workers implement a simple ABI to wrap their underlying decoder
- Each decoder's worker puts its result for an input into an output slot; inputs are pruned once all workers have attempted it
- Outputs are collected into "cohorts" of N for N decoders for a single input

## Mishegos: analysis framework

#### Fuzzing with mishegos produces cohorts of outputs

- Each output in a cohort contains the decoder state for the cohort's input:
  - Status (success, failure, kind of failure)
  - # of bytes decoded
  - Disassembly of instruction, length of disassembly

#### Cohorts need to be analyzed for errors and discrepancies

- Observation: analyses compose well, perform cheap + effective ones first
- Remove cohorts that only contain errors, only contain successes, ...
- Select cohorts where outputs disagree on status, length
- o Treat particular outputs as "high-quality" and use them as ground truth

## Mishegos: the bird's eye view

input	worker			
	./src/worker/bfd/bfd.so	./src/worker/capstone/capstone.so	./src/worker/xed/xed.so	./src/worker/zydis/zydis.so
6567f2f39cb2a58654	gs addr32 repnz repz pushf (5 / 28)	pushfq (5 / 8)	addr32 pushfq (5 / 14)	pushfq (5 / 6)
26f267664d0f3817314aecdf9d	es addr32 data16 rex.WRB (bad) (8 / 32)	ptest xmm14, xmmword ptr es:[r9d] (9 / 34)	(0 / 0)	(0 / 0)
656636f26f0f3a6f959066b1fd8c52	gs repnz outs dx,WORD PTR ss:[rsi] (5 / 35)	repne outsd dx, dword ptr ss:[rsi] (5 / 35)	repne data16 outsw gs (5 / 21)	repne outsw (5 / 11)
67f03ef0460f1104fe	lock lock movups XMMWORD PTR ds:[esi+r15d*8],xmm8 (9 / 50)	(0 / 0)	(0 / 0)	(0 / 0)
652e26520ffda71fd5bc9e3090235f	gs cs es push rdx (4 / 18)	push rdx (4 / 9)	push rdx (4 / 8)	push rdx (4 / 8)
64652e26520ffda71fd5bc9e309023	fs gs cs es push rdx (5 / 21)	push rdx (5 / 9)	push rdx (5 / 8)	push rdx (5 / 8)
6636f26f0f3a6f959066b1fd8c523e	repnz outs dx,WORD PTR ss:[rsi] (4 / 32)	repne outsd dx, dword ptr ss:[rsi] (4 / 35)	repne data16 outsw (4 / 19)	repne outsw (4 / 11)
2e26520ffda71fd5bc9e3090235f0d	cs es push rdx (3 / 15)	push rdx (3 / 9)	push rdx (3 / 8)	push rdx (3 / 8)
36f2f3f00f3a6315cd	ss repnz lock (bad) (7 / 21)	(0 / 0)	(0 / 0)	(0 / 0)
262ef0f00f3ada4eb5bae8	es cs lock lock (bad) (7 / 23)	(0 / 0)	(0 / 0)	(0 / 0)
f365f33e9559dd95a732b088057275	repz gs repz ds xchg ebp,eax (5 / 29)	xchg eax, ebp (5 / 14)	xchg ebp, eax (5 / 13)	xchg ebp, eax (5 / 13)
65f33e9559dd95a732b08805727509	gs repz ds xchg ebp,eax (4 / 24)	xchg eax, ebp (4 / 14)	xchg ebp, eax (4 / 13)	xchg ebp, eax (4 / 13)
f33e9559dd95a732b08805727509f3	repz ds xchg ebp,eax (3 / 21)	xchg eax, ebp (3 / 14)	xchg ebp, eax (3 / 13)	xchg ebp, eax (3 / 13)
3e9559dd95a732b08805727509f395	ds xchg ebp,eax (2 / 16)	xchg eax, ebp (2 / 14)	xchg ebp, eax (2 / 13)	xchg ebp, eax (2 / 13)
6765676547be4b69	addr32 (1 / 7)	(0 / 0)	(0 / 0)	(0 / 0)
59dd95a732b08805727509f39507d0	pop rcx (1 / 11)	pop rcx (1 / 8)	pop rcx (1 / 7)	pop rcx (1 / 7)
3e3e2e265ddf3b	ds ds cs es pop rbp (5 / 20)	pop rbp (5 / 8)	pop rbp (5 / 7)	pop rbp (5 / 7)
9559dd95a732b08805727509f39507	xchg ebp,eax (1 / 15)	xchg eax, ebp (1 / 14)	xchg ebp, eax (1 / 13)	xchg ebp, eax (1 / 13)
646526f04991a85e	fs gs es lock xchg r9,rax (6 / 26)	(0 / 0)	(0 / 0)	(0 / 0)
f26426644a0f38d5fbbe	repnz fs es fs rex.WX (bad) (8 / 29)	(0 / 0)	(0 / 0)	(0 / 0)
64666566950f3a13f9f6	fs data16 gs xchg bp,ax (5 / 24)	xchg ax, bp (5 / 12)	xchg bp, ax (5 / 11)	xchg bp, ax (5 / 11)
f02ef065440f5a52e209f49611d758	lock cs lock cvtps2pd xmm10,QWORD PTR gs:[rdx-0x1e] (9 / 52)	(0 / 0)	(0 / 0)	(0 / 0)
36266467470f3875ae022de4a1517e	ss es fs addr32 rex.RXB (bad) (8 / 31)	(0 / 0)	(0 / 0)	(0 / 0)
266467470f3875ae022de4a1517e90	es fs addr32 rex.RXB (bad) (7 / 28)	(0 / 0)	(0 / 0)	(0 / 0)
6467470f3875ae022de4a1517e90b4	fs addr32 rex.RXB (bad) (6 / 25)	(0 / 0)	(0 / 0)	(0 / 0)
f3f32666cd0f38b0c9a1ef83f720	repz repz es data16 int 0xf (6 / 28)	int 0xf (6 / 8)	int 0xf (6 / 7)	int 0x0F (6 / 8)
2ef066f0480f3a2904b7	cs lock data16 lock rex.W (bad) (8 / 33)	(0 / 0)	(0 / 0)	(0 / 0)
67470f3875ae022de4a1517e90b4cb	addr32 rex.RXB (bad) (5 / 22)	(0 / 0)	(0 / 0)	(0 / 0)
3664970f7a50db978a650a8288bee1	ss fs xchg edi,eax (3 / 19)	xchg eax, edi (3 / 14)	xchg edi, eax (3 / 13)	xchg edi, eax (3 / 13)
f226f236458a49fb848d28	repnz es repnz mov r9b,BYTE PTR ss:[r9-0x5] (8 / 44)	mov r9b, byte ptr ss:[r9 - 5] (8 / 30)	mov r9b, byte ptr [r9-0x5] (8 / 26)	mov r9b, ss:[r9-0x05] (8 / 21)

7

## Mishegos: performance, results, evaluation

- Fuzzer: 33M cohorts/hour → 228M decoder results/hour
- Analyzer:
  - **68.4M decoder results/hour** after deduplication
  - 130M net cohorts for a 4 hour campaign:
    - **15M status discrepancies** (one or more decoders disagree on validity)
    - **3.4M size discrepancies** (one or more decoders disagree on decoded length)
    - **59K cases** of XED "overaccepting" instructions
      - Probably all others underaccepting e.g. multi-byte NOPs!

Decoders: XED, libopcodes, Zydis, Capstone, DynamoRIO, bddisasm, Iced

Evaluated on 8 physical cores of an Intel Xeon 6140, running 7 decoder workers + the mutation engine and collector process

## Mishegos: analysis highlights

Discovered a variety of decoder bugs through discrepancies:

- Capstone: Incorrect control flow targets (call), incorrect length decoding esp. with legacy prefixes
- XED: False positives and false negatives for ISA extensions
- bddisasm: False positives, particularly around decoding instructions that don't work on AMD64 in long mode
- libopcodes (GNU bfd): Remarkably broken! Failed to decode a large number of valid instructions

#### Future research

Mishegos constitutes the first step (discovery) in evaluating the security posture of instruction decoders

#### Important future work:

- Evaluating the salience of discrepancies (control/data flow, decoder confusion, misaligning the instruction stream)
- Automatically emplacing discrepancies (generating "schizophrenic" binaries)
- Further refinement of "sliding"; minimizing duplicate work

#### Resources

GitHub: <a href="https://github.com/trailofbits/mishegos">https://github.com/trailofbits/mishegos</a>

Blog post: <u>Destroying x86-64 decoders with differential fuzzing</u>

Selected citations:

Paleari et al. 2010: N-version disassembly

Jay & Miller 2018: FLEECE

Domas 2017: Breaking the x86 ISA