




Differential fuzzing, or: how to find bugs when (ground) truth isn't real

William Woodruff
OSIRIS Lab

- **Me**

- *senior* security engineer (certified good at computers)
- R&D: program analysis research (mostly LLVM)
- Engineering: open-source C, C++, Rust, Ruby, Python




- **Trail of Bits**

- About 60 people, ~30% in NYC, rest remote
- Research, engineering, assurance
- Very good
- Summer and winter internships



Today's agenda

- Finding bugs
- Finding bugs with fuzzing
- Normal fuzzing can't find some bugs :(
 - We don't even know what bugs are, actually
- Spicy (differential) fuzzing to find those bugs
- Example case: x86_64 decoding
 - Demo!!!



Finding bugs

- Why? We want to...

- write reliable, safe code (cred)
- embarrass our coworkers (more cred)
- embarrass help our clients (money)


- How?

- Manual code review
 - Expensive (money & time), fallible (goto fail)
- Static analysis, formal methods, symbolic execution
 - Cheap-ish (compute heavy), sometimes effective, often infeasible (luv 2 explode state)
- Fuzzing

```
static OSStatus SSLVerifySignedServerKeyExchange(SSLContext *ctx,
bool isRsa, SSLBuffer signedParams, uint8_t *signature, UInt16 signatureLen)
{
    ...
    if ((err = ReadyHash(&SSLHashSHA1, &hashCtx)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &clientRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.verify(&hashCtx, &hashOut, signature, signatureLen)) != 0)
        goto fail;

    err = ...
    if(err) {
        sslErrorLog("SSLDecodeSignedServerKeyExchange: sslRawVerify "
                    "returned %d\n", (int)err);
        goto fail;
    }

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}
```



Finding bugs with fuzzing

- The TL;DR of fuzzing: feed garbage into the program until it crashes

```
while true; do  
    crash-monitor ./program -- /dev/urandom  
done
```

- Inexpensive, *shockingly* good at discovering (exploitable) bugs
- Problem: random inputs won't explore the program much
 - Intuition: Most random inputs don't resemble HTML, ZIP streams, PNGs...
- Solution: Use a feedback mechanism to guide inputs
 - "If it <runs longer/calls more functions/has more coverage>, try similar inputs"

EVOLUTION

GOD
GHT

SCIENCE IS A **LIAR** SOMET

ARISTOTLE



GALILEO



NEBULOSA ORION



NEBULOSA



Cool
LIGHT

Not all bugs are easily fuzzable

- Not all bugs...
 - cause easy-to-observe crashes (segfaults, aborts, non-zero exits, &c)
 - are memory corruptions (logic errors, permission errors, DoS, &c)
- If we have a specification, we can instrument the program to turn non-crashing errors into discoverable crashes

```
int x = get_some_untrusted_input();  
do_something_dangerous(x);
```



```
int x = get_some_untrusted_input();  
if (__is_invalid(x)) __crash();  
do_something_dangerous(x);
```

What's a bug without a specification?

- Not everything has a real specification
- Lots of things have “specifications” that are basically ignored
- The real spec is generally-agreed-upon behavior
 - “What does Adobe Acrobat do? Make our program do that”
- Lots of things are written in memory safe languages
 - == no memory corruption == no crashes on “bugs”
 - + no specification == no easy instrumentation approach :(

Another perspective: ground truth and oracles




- **Restate the problem: instead of bugs, we want an oracle**
- **Oracle supplies some notion of “ground truth”**
 - a yes or no answer for whether some behavior is correct
- **Different oracles:**
 - C and C++: segfaults, assertions, non-zero exits
 - Memory safe languages: exceptions, assertions, contract violations
- **Still no oracle if a “bug” doesn’t cause any of these!**
 - Back where we started :(

Constructing ground truth from difference

- **Observation: lots of things have multiple implementations**
 - Multiple PDF parsers, ZIP extractors, HTTP header parsers
- **Observation: lots of programs copy ideas and features from competitors**
 - “Acrobat can do \$X so our program needs to be able to do \$X!”
- **Observation: copying features without a specification means underspecification + lots of variation on unexpected inputs**
- **What if we compared different implementations?**
 - What if we define “bug” == “difference between impls”?



Differential fuzzing: we can't all be right



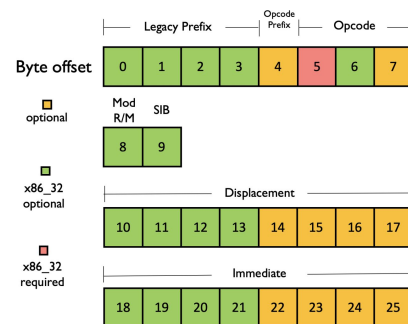
- Three programs, two different results (and , italic vs. bold)
- Not clear which is “right”, but both *probably* aren't
- No crashes needed!
- What if this but automated?

Differential fuzzing: applications

- **Any complex, popular format with competing implementations**
 - PDF, Word, media container formats (MKV, MP4)
- **Crypto primitives (hashing, digital signatures)**
 - Prior work: Wycheproof (Google), CDF (Kudelski)
- **Competing hardware implementations of ISAs**
 - x86_64: sandsifter
- **Competing software decoders for ISAs**
 - ARM: MC-Hammer
 - x86_64: mishegos (us!)

Case study: x86_64 decoding

- **Ideal target for differential fuzzing:**
 - Large, messy ISA with thousands of unique instructions
 - Complex encoding format with >50 years of backwards compatibility
 - Variable-length instructions, unlike ARM! Up to 15 bytes!!!
 - Two major vendors with totally independent implementations: Intel, AMD
 - Lots of popular, open source decoders to compare:
 - Capstone (mostly LLVM), zydis, XED (Intel), libopcodes (GNU)
 - High-interest/impact bugs:
 - Mess up debuggers, RE platforms, static analysis tools, ...

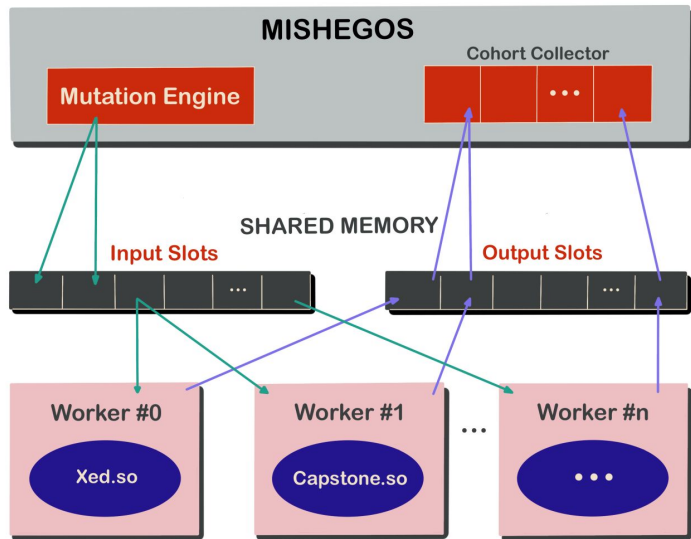


Case study: x86_64 decoding

The basic idea:

- Spawn a bunch of workers that wrap different decoder impls.
- Blast “random” inputs at the workers
 - Not really random: use x86_64’s structure to inform our choices
 - Legacy prefixes, SIB byte, &c
- Record what each worker claims each input decodes to
- Compare and contrast
- ?? Bugs ??

Mishegos: differential fuzzing for x86 decoders




input	worker	/src/worker/bfd/bfd.so	/src/worker/capstone/capstone.so	/src/worker/xed/xed.so	/src/worker/zydis/zydis.so
f02e2ef3440f3ac9	lock es cs repz rex.R (bad) (8 / 27)	(0 / 0)	(0 / 0)	(0 / 0)	(0 / 0)
f0f32664454b7002	lock repz es fs rex.RB (5 / 22)	jo 0xa (8 / 6)	(0 / 0)	(0 / 0)	(0 / 0)
36	ss (1 / 2)	(0 / 0)	(0 / 0)	(0 / 0)	(0 / 0)
3e2ef2f30f3aff	ds es repnz repz (bad) (7 / 22)	(0 / 0)	(0 / 0)	(0 / 0)	(0 / 0)
f0f326669d	lock repz es popfw (5 / 18)	popbt (5 / 4)	(0 / 0)	(0 / 0)	(0 / 0)
6764f3f30f50df	addr32 fs repz repz (bad) (6 / 25)	movmskps ebx, xmm7 (7 / 18)	(0 / 0)	(0 / 0)	(0 / 0)
f026f3670fa1	lock es repz addr32 pop fs (6 / 26)	pop fs (6 / 6)	(0 / 0)	(0 / 0)	(0 / 0)
f2	repnz (1 / 5)	(0 / 0)	(0 / 0)	(0 / 0)	(0 / 0)

Mishegos: making sense of the noise

- **Results: ~tens of millions of results per hour**
 - Depends on the number of workers, system load, ...
- **Need an automated strategy for filtering the interesting results**
- **Observation: we want a list (or DAG?) of filters to run, biggest first**
- **Implemented as “passes” (think LLVM) on transformed (JSON) output**
- **Boils down to this pipeline:**

```
mishegos workers.spec | mish2jsonl | analysis -p some-pass | mishmat
```

Contact Slide




William Woodruff

Senior Security Engineer

william@trailofbits.com

[@8x5clPW2](https://twitter.com/8x5clPW2) | github.com/woodruffw



TRAIL *OF* **BITS**